

CENG 3420

Computer Organization & Design



Lecture 06: Control Instruction

Bei Yu

CSE Department, CUHK

byu@cse.cuhk.edu.hk

(Latest update: January 27, 2022)

Spring 2022



- 1 Introduction
- 2 Control Instructions
- 3 Others
- 4 Summary



- 1 Introduction
- 2 Control Instructions
- 3 Others
- 4 Summary

RISC-V Instruction Fields



RISC-V fields are given names to make them easier to refer to

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0																	
funct7					rs2					rs1					funct3					rd					opcode					R-type	add
imm[11:0]										rs1					funct3					rd					opcode					I-type	addi
imm[11:5]					rs2					rs1					funct3					imm[4:0]					opcode					S-type	}
imm[12]		imm[10:5]			rs2					rs1					funct3					imm[4:1]		imm[11]			opcode					B-type	
imm[31:12]										rd										opcode					U-type	}					
imm[20]		imm[10:1]				imm[11]			imm[19:12]					rd					opcode					J-type							

B-format: no imm[0]

opcode 6-bits, opcode that specifies the operation

rs1 5-bits, register file address of the first source operand

rs2 5-bits, register file address of the second source operand

rd 5-bits, register file address of the result's destination

imm 12-bits / 20-bits, immediate number field

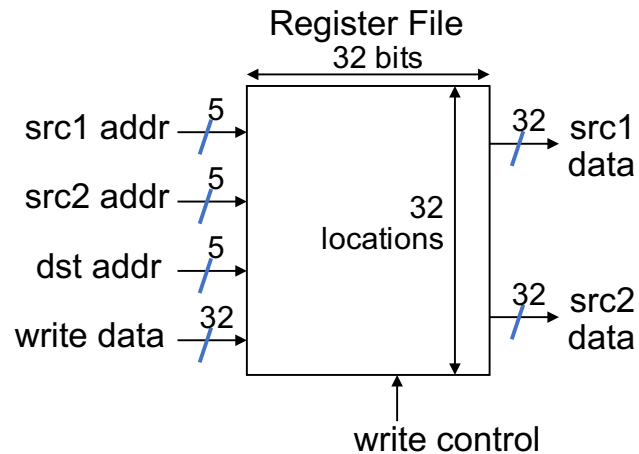
funct 3-bits / 10-bits, function code augmenting the opcode



Instruction Categories

- Load and Store instructions
- Bitwise instructions
- Arithmetic instructions
- Control transfer instructions
- Pseudo instructions

RISC-V Register File



- Holds thirty-two 32-bit general purpose registers
- Two read ports
- One write port

Registers are

- **Faster** than main memory
 - But register files with more locations are slower
 - E.g., a 64 word file may be 50% slower than a 32 word file
 - Read/write port increase impacts speed quadratically
- **Easier** for a compiler to use
 - $(A * B) - (C * D) - (E * F)$ can do multiplies in any order vs. stack
- Can hold variables so that code density improves (since register are named with fewer bits than a memory location)


Aside: RISC-V Register Convention



Table: Register names and descriptions

Register Names	ABI Names	Description
x0	zero	Hard-wired zero
x1	ra	Return address
x2	sp	Stack pointer
x3	gp	Global pointer
x4	tp	Thread pointer
x5	t0	Temporary / Alternate link register
x6-7	t1 - t2	Temporary register
x8	s0 / fp	Saved register / Frame pointer
x9	s1	Saved register
x10-11	a0-a1	Function argument / Return value registers
x12-17	a2-a7	Function argument registers
x18-27	s2-s11	Saved registers
x28-31	t3-t6	Temporary registers

History of RISC-V





- 1 Introduction
- 2 Control Instructions
- 3 Others
- 4 Summary

RISC-V Control Flow Instructions



RISC-V conditional branch instructions:

```
bne s0, s1, Lbl    # go to Lbl if s0 != s1
beq s0, s1, Lbl    # go to Lbl if s0 = s1
```

Example

Text

```
C: if (i==j) h = i + j;
```

```
    bne s0, s1, Lbl1    (assume s0=i; s1=j; s3=h)
    add s3, s0, s1
```

```
Lbl1:    ...
```

- Instruction Format (**B** format) B-format: no imm[0]
- How is the branch destination address specified ?

Q: why we don't need imm[0]? note imm is the target address.

— imm[0] is always zero, because address is word aligned. (for 32-bit cpu, one word contains 4 bytes)

In Support of Branch Instructions



- We have `beq`, `bne`, but what about other kinds of branches (e.g., `branch-if-less-than`)?
- For this, we need yet another instruction, `slt` (less than set)

Set on less than instruction:

```
slt t0, s0, s1      # if s0 < s1 then
                    # t0 = 1      else
                    # t0 = 0
```

- Instruction format (R format or I format)
`slti`

Alternate versions of `slt`

```
slti  t0, s0, 25    # if s0 < 25 then t0 = 1 ...
sltu  t0, s0, s1      # if s0 < s1 then t0 = 1 ...
sltiu t0, s0, 25      # if s0 < 25 then t0 = 1 ...
```

`sltu`: `slt` instruction for unsigned values

Aside: More Branch Instructions



Can use `slt`, `beq`, `bne`, and the fixed value of 0 in register `zero` to create other conditions

- less than: `blt s1, s2, Label`

`slt` `t0, s1, s2` *# t0 set to 1 if*
`bne` `t0, zero, Label` *# s1 < ~~s~~s2*

- less than or equal to: `ble s1, s2, Label`
- greater than: `bgt s1, s2, Label`
- great than or equal to: `bge s1, s2, Label`
- Such branches are included in the instruction set as **pseudo** instructions – recognized (and expanded) by the assembler

Bounds Check Shortcut



- Treating signed numbers as if they were unsigned gives a low cost way of checking if $0 \leq x < y$ (index out of bounds for arrays)

```
sltu t0, s1, t2      # t0 = 0 if
                      # s1 > t2 (max) (we know t2 is larger than 0)
                      # or s1 < 0 (min)
beq  t0, zero, IOOB  # go to IOOB if
                      # t0 = 0
```

- The key is that negative integers in two's complement look like large numbers in unsigned notation.
- Thus, an unsigned comparison of $x < y$ also checks if x is negative as well as if x is less than y .


Other Control Flow Instructions



- RISC-V also has an unconditional branch instruction or **jump** instruction:

```
jal zero, label           # go to label, label can be an  
                           immediate value
```

- Instruction Format (**J Format**)
- J is a pseudo instruction of unconditional jal and it will discard the return address (e.g., j label)



$pc := pc + \text{sign_extended}(\text{imm20} \ll 1)$



EX-2: Branching Far Away

What if the branch destination is further away than can be captured in 12 bits? Re-write the following codes.

```
beq    s0, s1, L1
```

Solution:



EX-2: Branching Far Away

What if the branch destination is further away than can be captured in 12 bits? Re-write the following codes.

```
beq    s0, s1, L1
```

note: beq the imm range is small, so we can use J instruction, if target is further away.

Solution:

```
    bne s0, s1, L2
    j  L1
L2:  ...
```

j is a pseudo instruction for jal (i.e., jal zero, label)



EX: Compiling a while Loop in C

```
while (save[i] == k) i += 1;
```

Assume that `i` and `k` correspond to registers `s3` and `s5` and the base of the array `save` is in `s6`.



EX: Compiling a while Loop in C

```
while (save[i] == k) i += 1;
```

Assume that *i* and *k* correspond to registers *s3* and *s5* and the base of the array *save* is in *s6*.

```
                                (left shift 2 == x 4)
Loop: slli                        # Temp reg t1 = i * 4    t1: address shift in array
      sll  t1, s3, 2              # t1 = address of save[i]
      add  t1, t1, s6             #
      lw   t0, 0(t1)             # Temp reg t0 = save[i]
      bne  t0, s5, Exit          # go to Exit if save[i] != k
      addi s3, s3, 1             # i = i + 1
      j    Loop                 # j is a pseudo instruction for jal
                                # go to Loop

Exit:
```

Note: left shift *s3* to align word address, and later address is increased by 1

32-bit CPU: one word == 32bits == 4 bytes

64-bit CPU: one word == 64bits == 8 bytes

Six Steps in Execution of a Procedure



- ① Main routine (**caller**) places parameters in a place where the procedure (**callee**) can access them
 - a0 – a7: four argument registers
- ② **Caller** transfers control to the **callee**
- ③ **Callee** acquires the storage resources needed
- ④ **Callee** performs the desired task
- ⑤ **Callee** places the result value in a place where the **caller** can access it
 - s0–s11: 12 value registers for result values
- ⑥ **Callee** returns control to the **caller**
 - **ra**: one return address register to return to the point of origin



Instructions for Accessing Procedures



We have learnt `jal`, now let's continue

- RISC-V procedure call instruction:

```
jal ra, label # jump and link,  
           # label can be an immediate value
```

calling address:  `add xxx`
return address:  `jal ra, label`
...
...

label:

- Saves **PC + 4** in register `ra` to have a link to the next instruction for the procedure return
- Machine format (**J** format):
- Then can do procedure return with a

PC: point counter. current instruction location.
when we move to next instruction, PC+4

```
jalr x0, 0(ra) # return
```

- Instruction format (**I** format)

Example of Accessing Procedures



- For a procedure that computes the GCD of two values i (in $t0$) and j (in $t1$):
`gcd(i, j);`
- The caller puts the i and j (the parameters values) in $a0$ and $a1$ and issues a

```
jal (ra), gcd      # jump to routine gcd
```

- The callee computes the GCD, puts the result in $s0$, and returns control to the caller using

```
gcd: . . .          # code to compute gcd  
     jalr x0, 0(ra)    # return
```




What if the callee needs to use **more registers** than allocated to argument and return values?

- Use a **stack**: a last-in-first-out queue
- One of the general registers, **sp**, is used to address the stack
- “grows” from high address to low address
- **push**: add data onto the stack, data on stack at new sp

$$sp = sp - 4$$

- **pop**: remove data from the stack, data from stack at sp

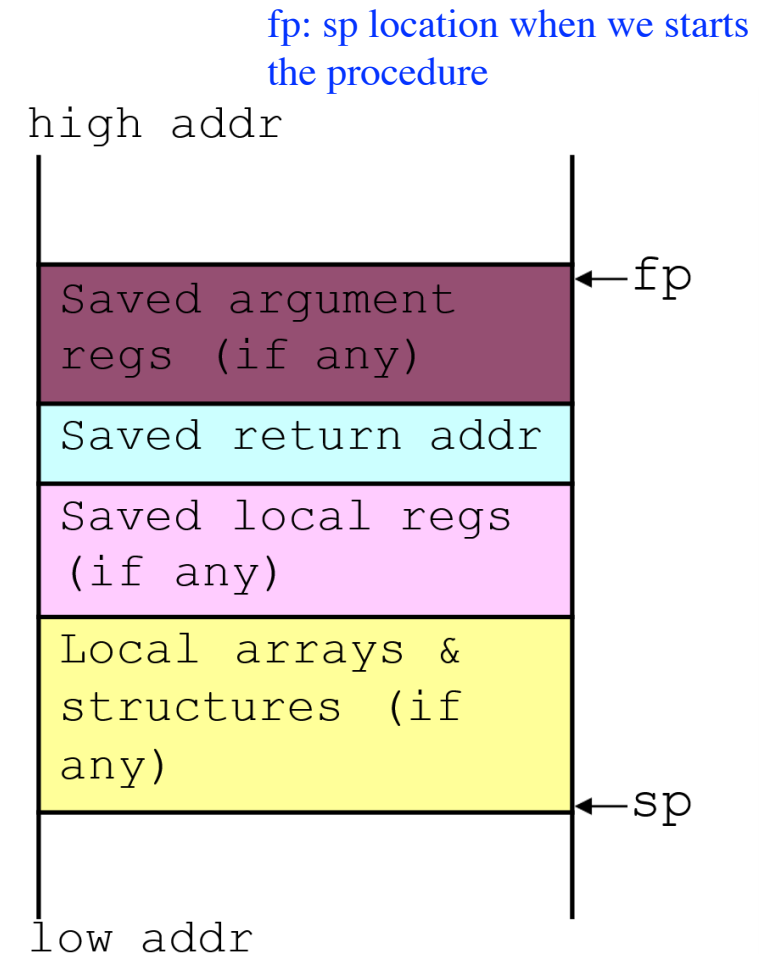
$$sp = sp + 4$$



Allocating Space on the Stack



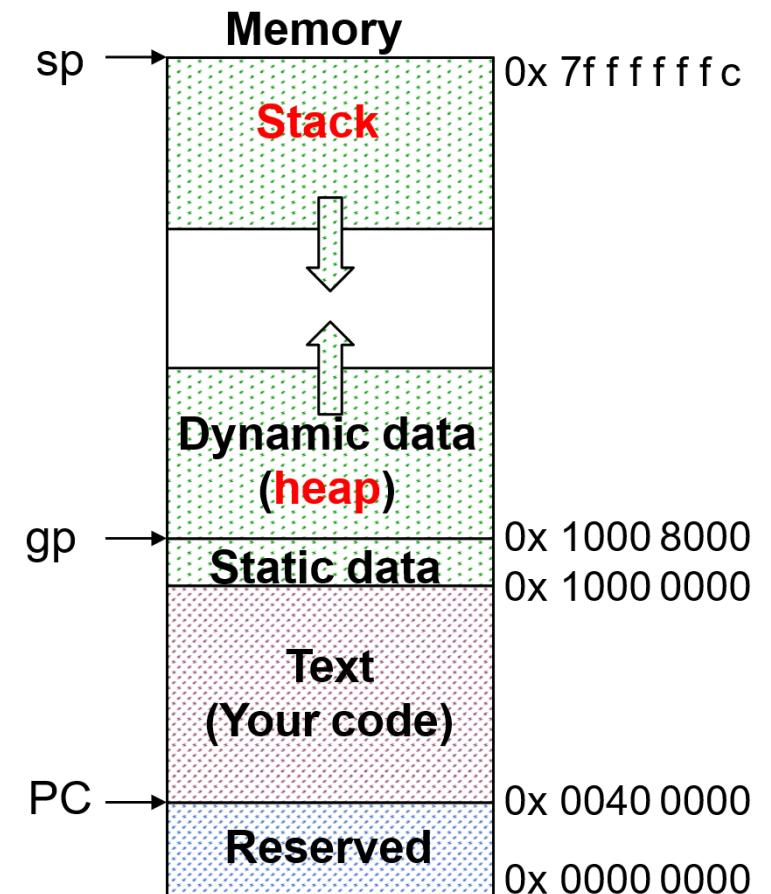
- The segment of the stack containing a procedure's saved registers and local variables is its procedure frame (aka activation record)
- The frame pointer (**fp**) points to the first word of the frame of a procedure – providing a stable “base” register for the procedure
- fp is initialized using sp on a call and sp is restored using fp on a return



Allocating Space on the **Heap**



- Static data segment for constants and other static variables (e.g., arrays)
- Dynamic data segment (aka heap) for structures that grow and shrink (e.g., linked lists)
- Allocate space on the heap with `malloc()` and free it with `free()` in C





EX-3: Compiling a C Leaf Procedure

Leaf procedures are ones that do not call other procedures. Give the RISC-V assembler code for the follows.

```
int leaf_ex (int g, int h, int i, int j)
{
    int f;
    f = (g+h) - (i+j);
    return f;
}
```

Solution:



EX-3: Compiling a C Leaf Procedure

Leaf procedures are ones that do not call other procedures. Give the RISC-V assembler code for the follows.

```
int leaf_ex (int g, int h, int i, int j)
{
    int f;
    f = (g+h) - (i+j);
    return f;
}
```

Solution:

Suppose g, h, i, and j are in a0, a1, a2, a3

```
leaf_ex:  addi    sp, sp, -8 # make stack room (we push two values into the stack)
```

```
        { sw      t1, 4(sp) # save t1 on stack
```

```
        { sw      t0, 0(sp) # save t0 on stack
```

```
        add      t0, a0, a1
```

```
        add      t1, a2, a3
```

```
        sub      s0, t0, t1
```

```
        { lw      t0, 0(sp) # restore t0
```

```
        { lw      t1, 4(sp) # restore t1
```

```
        addi     sp, sp, 8 # adjust stack ptr
```

```
        jalr     zero, 0(ra)
```

Return to caller

contains address of caller

Q: why we need to store t0 & t1 into the stack?
—other procedures may also use t0 & t1;
since locally we may use t0 & t1, we need to backup
their original values